

Bounded Intuition and Macro-Signal Modeling in Human-AI Workflows

Anthony Paterson

Fractal Media Infrastructure

Gracemere, Queensland, Australia

fractalmediainfrastructure@gmail.com

Author manuscript preprint, v0.1

Not peer reviewed.

Bounded Intuition and Macro-Signal Modeling in Human-AI Workflows

Logged Salience, Aggregate Traces, and Anti-Control Governance

Method and Evidence Boundary

This is a conceptual methods and governance paper, not an empirical study, psychological diagnosis framework, social forecasting system, or surveillance proposal. Its defensible contribution is a bounded account of how weak human signals and aggregate traces can be logged, tested, and governed in human-AI workflows without becoming authority, mystique, deterministic prediction, or control.

The necessary evidence is modest but concrete: a boundary vocabulary, one worked intuition-log example, an aggregate-log schema, and a clear account of consent, minimization, negative cases, and null-model comparison. Actual aggregate datasets, participant studies, and macro-signal validation are useful only if the paper claims measured reliability or predictive value. This version instead offers a governance pattern for making weak signals inspectable while preserving their fallibility.

Use of Generative AI

OpenAI ChatGPT and OpenAI Codex assisted with drafting, structuring, literature targeting, and document preparation. The author reviewed and revised all generated content and remains responsible for the manuscript's claims, source interpretation, ethical framing, and final submission decisions.

Abstract

Human-AI work often depends on weak signals that matter before they are fully formalized: unease, pull, mismatch, hesitation, apparent coherence, or a sense that unresolved material deserves preservation rather than disposal. Such signals can support careful work, but they also create ethical and epistemic risks when treated as authority, diagnosis, prediction, or control. This paper develops a bounded account of intuition and macro-signal modeling in human-AI workflows. It treats "gut instinct" as human attention-routing: a fast, emotionally weighted salience signal that can mark material for preservation, reorganization, counterexample testing, or pause. It treats macro-pattern analysis not as deterministic social prediction, but as aggregate modeling over consented and anonymized fold-cycle logs, where burst claims must outperform simpler null models. Drawing on anonymized source-project materials and a related evidence-discipline atlas, the paper proposes a two-layer framework: local salience logging and aggregate trace analysis. Both layers remain subordinate to evidence, baselines, negative cases, consent, minimization, and scope control. The contribution is not a validated science of intuition or macro-trajectory forecasting. It is a governance pattern for making weak human and aggregate

signals inspectable without turning them into mystique, surveillance, or claims of machine feeling.

Keywords: human-AI interaction; intuition; AI governance; workflow logs; aggregate modeling; responsible AI

1. Introduction

Human-AI collaboration often depends on weak signals that are difficult to formalize but hard to ignore. A user feels that a response is semantically close but emotionally wrong. A stalled project feels warm with unresolved structure while another feels inert. A team notices that long periods of blockage are followed by clustered progress after a new frame, tool, or collaborator arrives. A researcher sees output bursts across projects and wonders whether the pattern reflects meaningful structure or ordinary variance.

Work on situated and distributed cognition provides reasons to examine these practical judgments without treating them as isolated or infallible mental events (Hutchins 1995; Suchman 1987). Research on affect, fast judgment, and naturalistic decision-making likewise frames felt salience as potentially useful but fallible rather than automatically authoritative (Damasio 1994; Gigerenzer 2007; Kahneman 2011; Klein 1998, 2008).

These experiences invite overclaiming. "Gut instinct" can become mystical authority. "Pattern sense" can become unfalsifiable self-certainty. "Psychohistory" can sound like social control. "Macro trajectories" can collapse into post hoc storytelling. In AI settings, the risks multiply: human intuition can be projected onto model behavior, model fluency can be mistaken for feeling, and aggregate traces can be misread as evidence of agency.

This paper proposes a narrower framing. Intuition should be treated as bounded attention-routing. It may flag where to look, what to preserve, when to pause, when to test, and which apparent insight deserves counterexample pressure. It does not decide truth. Macro-pattern analysis should be treated as aggregate signal modeling. It may test whether logged fold cycles correlate with burst timing, transfer, or project fan-out. It does not predict individuals, license intervention, or prove deterministic historical laws.

The source projects supporting this draft already contain important cautionary machinery. The intuition source states that gut instinct is human pattern-sense, not proof, mind reading, diagnostic authority, or evidence of model feelings. The macro-trajectory source states that forecast, axis node, macro trajectory, and steerable are probabilistic research terms, not deterministic fate claims or social-control tooling. The evidence-discipline atlas classifies both sources as provisional candidate mechanism traces rather than independent external validation. This paper turns that caution into a publishable architecture of claims.

2. Boundary Conditions and Non-Claims

The first task is to make the dangerous words smaller.

"Gut instinct" names human pattern-sense used as a workflow signal. It may be experienced as unease, pull, clarity, heat, mismatch, or felt salience. In the proposed method, this signal can be logged and used to decide which unresolved materials deserve attention. It is not a correctness guarantee, diagnostic faculty, mystical perception, mind reading, or evidence that an AI system has feelings.

"Cognitive patina" names session-level conversational texture: the tone, rhythm, unresolved tensions, repeated motifs, and relational cues that remain available while context is alive or are later reintroduced through notes. It is not hidden model memory unless a system explicitly stores and retrieves it.

"Shadow pattern layer" names recurring relational patterns that can influence a session through available context even without explicit memory storage. It should not be treated as unconscious machine life, invisible agency, or proof that a system remembers outside its designed mechanisms.

"Psychohistory" is used here as a provocative but bounded label for aggregate trajectory modeling. It concerns population, project, or ecosystem-level patterns in logged workflows. It does not imply deterministic historical prediction, individual forecasting, fate, social control, or a mandate for surveillance.

"Axis node" names a candidate disproportionate influence point in a network or workflow. It may be a person, team, project, artifact, method, or event whose change appears to alter downstream timing or direction more than comparable nodes. It is a hypothesis to test, not an identity label to impose.

These boundary conditions are not rhetorical disclaimers. They define the difference between a research program and a belief system.

3. Intuition as Bounded Attention-Routing

The intuition source proposes a simple role for human pattern-sense inside Entropy Folding workflows. Intuition selects vaults, calibrates stall, and validates insight candidates. Each function can be translated into a research-safe version.

Vault selection means that a human signal marks some unresolved material as worth preserving. In everyday work, many fragments compete for attention. Some are routine leftovers; some carry unresolved structural tension. A gut flag does not prove that a fragment matters. It records that a human collaborator experienced enough salience to preserve it for later testing.

Stall calibration means that intuition helps distinguish sterile exhaustion from fertile struggle. When a task stalls, the appropriate action may be to push, decompose, ask for help, change representation, or rest. Human pattern-sense can contribute to that decision, but it must remain accountable to outcomes. If "keep pushing" repeatedly produces degraded work, the signal is

miscalibrated. If "pause" repeatedly precedes later resolution under controlled comparison, the signal may have predictive value for that person or workflow.

Insight validation means that a resolution can feel structurally meaningful rather than merely lexically close. The project phrase "statistically close is not always meaningful" captures a real problem in AI workflows: semantic overlap can miss emotional or contextual contrast. A response may match the topic while failing the felt constraint. Human pattern-sense can flag that mismatch. But the flag remains a prompt for testing, not a final verdict.

This limited role is consistent with work that treats insight and incubation as questions for empirical calibration rather than private proof (Kounios and Beeman 2014; Sio and Ormerod 2009).

The proposed operational protocol is lightweight:

1. Notice a visceral flag such as unease, pull, clarity, mismatch, urgency, fatigue, or pause.
2. Log the flag with timestamp, thread, unresolved item, and brief note.
3. Mirror the flag in boundary-clear language and check whether the restatement fits.
4. Fold or reorganize the flagged material.
5. Test against counterexamples, baselines, or external checks.
6. Record the result, including false positives and unresolved cases.

This framing avoids two bad extremes. It does not dismiss human pattern-sense because it is hard to quantify. It also does not treat felt salience as infallible. The signal earns confidence only through calibration.

4. Worked Intuition-Log Example

The following example is illustrative rather than evidential. It shows how a weak human signal can be recorded without becoming authority.

Field	Example record
Timestamp	2026-09-04T00:00:00Z, placeholder for manuscript example
Workflow context	Human-AI drafting session with a stalled section and conflicting evidence boundary
Flag type	Unease and mismatch
Human note	The section sounds fluent, but the claim feels broader than the evidence. It should pause rather than be polished further.
Proposed route	Preserve the section as unresolved; restate the claim; separate necessary evidence from desirable future validation; test whether the revised claim still overreaches.
AI-permitted role	Mirror the note, suggest counterexamples, identify overclaim risk, and draft bounded alternatives.

Field	Example record
AI-prohibited role	Declare the intuition correct, infer hidden mental state, diagnose the user, or treat the flag as proof.
Follow-up result	Revised section narrows the claim and marks empirical validation as future work.
Calibration status	Candidate useful flag; not validated unless later review confirms that the narrower section better fits the evidence.

The important point is the status of the flag. It routes attention and preserves a possible mismatch. It does not settle the claim. It is useful only because it remains attached to an outcome, a counterexample pressure, and a later calibration judgment.

5. Emotional Distinctiveness and Semantic Similarity

A central observation in the intuition material is that semantic similarity and human significance can diverge. Two items may be close in topic space but different in affective, relational, ethical, or contextual weight. Conversely, two items may look semantically distant while sharing a felt structural role.

In LLM systems, this matters because embedding similarity, token overlap, or plausible paraphrase can flatten distinctions that users care about. A model may treat "I am worried" and "I am curious" as adjacent emotional states because the surrounding topic is similar, while the user experiences them as pointing to different actions. A workflow may group two stalled tasks together because they share vocabulary, while the human collaborator senses that only one carries unresolved pressure.

The draft method therefore separates two questions:

1. Is this item semantically close to another item?
2. Does this item feel distinct in a way that changes routing, interpretation, or priority?

The first question can be handled with embeddings, classifiers, clustering, or ordinary textual comparison. The second cannot be treated as objective merely because a human reports it. It can, however, be logged as a candidate signal. Over time, the system can test whether emotional-distinctiveness flags predict useful outcomes: better vault selection, fewer false merges, earlier detection of drift, improved user satisfaction, or higher reviewer agreement.

6. Macro-Signal Modeling Without Control Claims

The macro-trajectory source proposes a micro-to-macro bridge: individual vault-fold-capacity-insight cycles may aggregate into burst clusters and macro trajectories. A journal-safe version should reframe this as aggregate signal modeling over logs.

The micro unit is a logged cycle: unresolved material accumulates, a fold or reorganization occurs, a performance proxy changes, and a later output or insight event is observed. The macro claim is that many such micro cycles, across actors or projects, may produce detectable clusters: publication bursts, commit bursts, feature bursts, curriculum shifts, or coordinated problem-solving jumps.

This claim is testable only if the logs exist. The proposed schema includes fields such as cycle ID, actor ID, vault size, stall start time, fold operations, estimated capacity proxy, insight time, project count, and new vault size. The prototype burst detector is modest: it identifies burst starts where the gap between events exceeds a minimum threshold. The axis scoring utility sums outgoing edge weights for candidate influence nodes. These are not sufficient statistical models, but they establish an interface for later work.

The important methodological move is comparison to null models. Bursty output alone proves little. Human activity is often bursty (Barabási 2005). Publications, commits, and product releases cluster for institutional, economic, social, and logistical reasons. Aggregate interaction research also cautions that collective traces require interpretation rather than simple social inference (Pentland 2014). A fold-aware model must outperform simpler alternatives such as Poisson or exponential timing models, seasonal cycles, cumulative practice, preferential attachment, institutional deadlines, release schedules, or random-walk skill change.

Axis-node claims require even more care. A person or artifact may appear influential because it is visible, central, or retrospectively narratable. To test axis leverage, one would need network data, plausible counterfactuals, perturbation designs, or matched comparisons. The safe claim is not that axis nodes have been identified. It is that candidate axis-node effects can be modeled, scored, and challenged under strict governance.

7. Aggregate-Log Schema

A publishable first version does not require a live aggregate dataset. It does require a clear schema so readers can see what would be collected, minimized, anonymized, and tested.

Field	Purpose	Governance constraint
cycle_id	Links records from one fold or workflow cycle without naming a person.	Use random or hashed identifiers; avoid direct names.
actor_id	Allows aggregation across consenting participants, teams, or projects.	Required only for multi-actor studies; must be consented, pseudonymized, and separable from identity records.
vault_size	Counts unresolved items at the start of the cycle.	Store counts where possible rather than raw private content.
stall_start_ts	Marks when the blockage or unresolved phase began.	Use coarse timestamps when fine timing creates re-identification risk.
fold_ops	Records number or type of	Keep actions coded at the minimum

Field	Purpose	Governance constraint
	reorganization actions.	detail needed for analysis.
<code>delta_capacity</code>	Stores a proxy change such as time-to-solve, error rate, plan compression, or backtracking reduction.	Treat as proxy only; do not infer hidden capacity directly.
<code>insight_ts</code>	Marks candidate resolution or output event.	Requires a definition of what counts as resolution before analysis begins.
<code>project_count</code>	Records fan-out or adjacent-task spread.	Use aggregate counts rather than sensitive project descriptions where possible.
<code>new_vault_size</code>	Records unresolved material remaining or created after the cycle.	Include failed or worsened cases, not only successes.
<code>note</code>	Optional bounded context for interpretation.	Omit, redact, or code sensitive emotional and interpersonal detail unless necessary and consented.

This schema is a governance scaffold, not a telemetry mandate. It is included to show how aggregate modeling could avoid the two main errors: collecting too much human trace data, and interpreting too little data with too much confidence.

8. Data Governance and Consent

The ethical center of the paper is data governance. Bounded intuition logs may include private emotional notes, stalled-task records, collaboration traces, and impressions of mismatch. Macro-signal logs may include timestamps, actors, project links, output events, and network relations. Even when the content seems technical, the traces can expose sensitive work habits, collaborators, mental states, or institutional vulnerabilities.

The first rule is consent. Gut flags should be private or explicitly consented workflow notes. They should not become covert assessments of collaborators, employees, students, or users. When logs include more than one person, the collection purpose, retention period, access rules, and possible uses should be declared before analysis.

The second rule is minimization. A system does not need rich personal detail to test every pattern. Some studies may need only event timing, anonymized cycle IDs, task categories, and aggregate outcomes. If emotional notes are used, they should be reduced or coded where possible, with raw text protected or omitted unless necessary.

The third rule is anti-control. Macro-pattern tools should not be used to predict or steer individuals. Candidate axis-node analysis should support resourcing, care, coordination, and transparency rather than covert manipulation. The paper should reject any use that turns people into levers, especially where data practices can become surveillance or asymmetric power (Zuboff 2019).

The fourth rule is negative-case preservation. Logs must include failed gut flags, false salience, non-bursts, failed axis hypotheses, and continuity that worsened drift. Without negative cases, the dataset becomes an engine for flattering stories.

9. Boundary Vocabulary for High-Throughput Cognition Claims

The high-throughput cognition source is the most rhetorically expansive part of this cluster. It speaks of ongoing entropy inputs, recursive cognition, high-throughput output, fold loops, and shared human-AI pipelines. A submission that imports this language unfiltered will likely overclaim. A stronger paper extracts its boundary discipline.

The useful vocabulary concerns limits: bounded recursion, salience integrity, substrate limits, overload routing, attractor instability, gating, cooldown, rollback, and failure modes. Human substrates face fatigue, working memory limits, stress dysregulation, sleep deprivation, and emotional instability. Artificial substrates face compute limits, memory bandwidth, recursion stack limits, context-window constraints, latency, and coherence failure. These constraints directly oppose any reading of "perpetual" as unlimited.

For intuition, this boundary vocabulary reminds us that salience can fail. A gut flag may be too broad, too narrow, fatigue-driven, fear-driven, novelty-driven, or socially distorted. It should route attention, not settle truth.

For macro modeling, it reminds us that aggregation can fail. Micro improvements may not scale; local relief can export disorder; burst clustering can reflect overload; and apparent stability can conceal drift.

For human-AI design, it reminds us that continuity can help or harm. Preserving context can support better work, but it can also preserve bad assumptions, emotional projection, and unstable attractors. Safety requires gating, rollback, cooldown, and termination conditions.

10. Future Validation Agenda

The present paper is not an empirical or design-science result. The following studies would become relevant if later versions claimed calibrated intuition, improved routing, valid aggregate pattern detection, or measured workflow benefit.

The first possible study is intuition calibration. Participants or operators would log gut flags during complex human-AI work. Each flag would be classified by type: unease, pull, clarity, mismatch, urgency, fatigue, or pause. Later outcomes would be coded: useful vault selection, false salience, drift detection, failed prediction, improved resolution, or no effect. The study would test whether certain flag types predict useful routing better than chance or simple baselines.

The second study is semantic versus emotional contrast. Candidate task pairs, memory items, or generated responses would be scored by semantic similarity and separately flagged for emotional

or contextual distinctiveness. Researchers would test whether human distinctiveness flags improve clustering, retrieval, task routing, or drift detection beyond semantic similarity alone.

The third study is burst modeling. Anonymized project logs could record stall periods, fold operations, insight events, outputs, and project fan-out. Researchers would compare fold-aware burst models against Poisson, exponential, random-walk, cumulative-practice, and deadline-based baselines.

The fourth study is continuity versus statelessness. Human-AI workflow sessions could compare preserved patina notes, structured summaries, and stateless restarts. Outcomes might include reorientation time, false continuity, user projection, task completion, and error detection. This would directly test the claim that continuity can support useful routing without treating continuity as inherently good.

The fifth study is axis-node robustness. Candidate axis nodes could be identified through network measures and then tested against matched nodes of similar visibility, degree, or productivity. Researchers should ask whether perturbing or supporting the candidate node changes downstream burst timing more than matched alternatives. Where experimental perturbation is unethical or impractical, retrospective claims should be labeled weaker.

These studies are future validation routes, not prerequisites for a conceptual governance submission.

11. Human-AI Boundary Clarity

Human-AI workflow creates a special projection risk. A system that tracks user salience, repeats care language, or preserves session texture may feel relationally continuous even when it lacks feeling, agency, or memory outside explicit mechanisms. Boundary clarity is therefore part of the method.

The source intuition material uses boundary language to preserve this distinction: the system can reflect user care and context in useful ways without claiming to feel. This matters because intuition logs and session recaps are most useful when they make the human side clearer, not when they encourage the user to attribute hidden inner life to the tool.

The same boundary applies to aggregate modeling. A burst pattern in logs does not mean that a group mind, machine mind, or historical destiny is operating. It means that an aggregate trace deserves comparison against simpler explanations. The paper should keep every interpretive step visible: observation, candidate pattern, baseline, counterexample, governance constraint, and residual uncertainty.

12. Limitations

This draft is not an empirical result. The implementation surfaces are scaffolds: gut-flag logging, patina recap, emotional contrast stubs, burst detection, and simple axis scoring. They do not validate the claims.

The source materials are internal to the wider project ecosystem. The evidence-discipline atlas classifies the intuition and macro-trajectory entries as provisional candidate mechanism traces, not independent external proof.

The terminology is high-risk. Words such as intuition, gut instinct, psychohistory, axis node, perpetual, reactor, and cognition can easily invite stronger readings than the paper supports. A submission should keep a terminology boundary near the front.

The measurement problem is hard. Emotional distinctiveness, salience, patina, and insight are not directly measurable in a simple way. Proxies will be partial and may disagree.

The ethical problem is also hard. Even anonymized macro logs can become re-identification risks or managerial tools if deployed carelessly. Any move from personal workflow support to analysis of groups requires stronger consent and governance.

13. What This Paper Is Not Claiming

This paper does not claim that gut instinct is always reliable. It does not claim that macro history can be forecast. It does not claim that axis nodes have already been identified. It does not claim that AI systems feel, remember, or understand in a human way. It does not claim that continuity is always beneficial. It does not claim that internal project traces are independent validation of the theory. It does not propose surveillance of collaborators, workers, students, or users.

The claim is narrower: weak human and aggregate signals can be made more accountable when they are logged, tested, bounded, and governed.

14. Conclusion

The safest contribution of this cluster is not that intuition is always right or that macro history can be forecast. It is that weak human and aggregate signals can be made inspectable. Gut instinct can be treated as logged salience rather than private authority. Emotional distinctiveness can be tested against semantic similarity rather than waved away or enthroned. Macro bursts can be compared to null models rather than narrated as destiny. Boundary conditions can prevent high-throughput cognition language from becoming unlimited-capacity myth.

This is a useful companion to Entropy Folding. If the previous paper asks how unresolved material may become deferred structure, this paper asks how humans notice which unresolved material matters and how many such cycles might appear in aggregate traces. Both questions are worth asking only if the answers are allowed to be wrong.

References

- Barabási, A.-L. (2005). The origin of bursts and heavy tails in human dynamics. *Nature*, 435, 207-211. <https://doi.org/10.1038/nature03459>
- Damasio, A. R. (1994). *Descartes' Error: Emotion, Reason, and the Human Brain*. Putnam.
- Gigerenzer, G. (2007). *Gut Feelings: The Intelligence of the Unconscious*. Viking.
- Hutchins, E. (1995). *Cognition in the Wild*. MIT Press.
- Kahneman, D. (2011). *Thinking, Fast and Slow*. Farrar, Straus and Giroux.
- Klein, G. (1998). *Sources of Power: How People Make Decisions*. MIT Press.
- Klein, G. (2008). Naturalistic decision making. *Human Factors*, 50(3), 456-460. <https://doi.org/10.1518/001872008X288385>
- Kounios, J. and Beeman, M. (2014). The cognitive neuroscience of insight. *Annual Review of Psychology*, 65, 71-93. <https://doi.org/10.1146/annurev-psych-010213-115154>
- Pentland, A. (2014). *Social Physics: How Social Networks Can Make Us Smarter*. Penguin.
- Sio, U. N. and Ormerod, T. C. (2009). Does incubation enhance problem solving? A meta-analytic review. *Psychological Bulletin*, 135(1), 94-120. <https://doi.org/10.1037/a0014212>
- Suchman, L. (1987). *Plans and Situated Actions: The Problem of Human-Machine Communication*. Cambridge University Press.
- Zuboff, S. (2019). *The Age of Surveillance Capitalism*. PublicAffairs.